

An Introduction to Multi-label Learning (ML-KNN & BP-MLL)

Bobby Lumpkin



1 Introduction to Multi-label Learning

- Overview and Advantages

2 ML-KNN Approach

- Model Outline
- Computing Model Probabilities
- Implementation (in scikit-multilearn)

3 BP-MLL Approach

- Feed-forward Neural Networks
- Neural Network Loss Functions & Training for MLL
- Implementation (in TensorFlow/Keras)

Introduction to Multi-label Learning

What is it?

What is it?

- Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

What is it?

- Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.
- Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.

What is it?

- Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.
- Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.
 - * **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: “News”, “Health” and “Government”.

Multi-label Learning

What is it?

- Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.
- Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.
 - * **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: “News”, “Health” and “Government”.

Why use novel approaches?

Multi-label Learning

What is it?

- Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.
- Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.
 - * **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: “News”, “Health” and “Government”.

Why use novel approaches?

- Naive Approach: Train a sequence of independent binary classifiers (one per category)

Multi-label Learning

What is it?

- Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.
- Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.
 - * **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: “News”, “Health” and “Government”.

Why use novel approaches?

- Naive Approach: Train a sequence of independent binary classifiers (one per category)
- Doesn't capitalize on the information in the correlations between the different labels of each instance.

Multi-label Paradigm: Definitions & Notation

- Let \mathcal{X} denote the domain of instances and $\mathcal{Y} = \{1, \dots, Q\}$ be the finite set of labels.
- Given $x \in \mathcal{X}$ and its associated $Y \subseteq \mathcal{Y}$, let \vec{y}_x be the category vector for x such that (for all $\ell \in \mathcal{Y}$) $\vec{y}_x(\ell) = 1$ if $\ell \in Y$. Otherwise, $\vec{y}_x(\ell) = 0$.

ML-KNN Approach

ML-KNN Algorithm: More Notation

Notation:

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x , identified in the training set.

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x , identified in the training set.
- Let $\vec{C}_x(\ell) = \sum_{a \in N(x)} \vec{y}_a(\ell)$ ($\ell \in \mathcal{Y}$) define a membership counting vector.

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x , identified in the training set.
- Let $\vec{C}_x(\ell) = \sum_{a \in N(x)} \vec{y}_a(\ell)$ ($\ell \in \mathcal{Y}$) define a membership counting vector.
- Let H_0^ℓ denote the event that test instance t does not have a label ℓ and let H_1^ℓ denote the event that it does have label ℓ .

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x , identified in the training set.
- Let $\vec{C}_x(\ell) = \sum_{a \in N(x)} \vec{y}_a(\ell)$ ($\ell \in \mathcal{Y}$) define a membership counting vector.
- Let H_0^ℓ denote the event that test instance t does not have a label ℓ and let H_1^ℓ denote the event that it does have label ℓ .
- Let E_j^ℓ ($j \in \{1, \dots, K\}$) denote the event that, among the K nearest neighbors of t , there are exactly j instances which have label ℓ .

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

- Given a test instance, t , \vec{Y}_t is determined using the MAP estimate:

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

- Given a test instance, t , \vec{Y}_t is determined using the MAP estimate:

$$\begin{aligned}\vec{y}_t(\ell) &= \operatorname{argmax}_{b \in \{0,1\}} \mathbb{P} \left(H_b^\ell | E_{\vec{C}_t(\ell)}^\ell \right), \quad \ell \in \mathcal{Y} \\ &= \operatorname{argmax}_{b \in \{0,1\}} \frac{\mathbb{P} \left(H_b^\ell \right) \cdot \mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell \right)}{\mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell \right)} \\ &= \operatorname{argmax}_{b \in \{0,1\}} \mathbb{P} \left(H_b^\ell \right) \cdot \mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell \right)\end{aligned}$$

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

- Given a test instance, t , \vec{Y}_t is determined using the MAP estimate:

$$\begin{aligned}\vec{y}_t(\ell) &= \operatorname{argmax}_{b \in \{0,1\}} \mathbb{P} \left(H_b^\ell | E_{\vec{C}_t(\ell)}^\ell \right), \quad \ell \in \mathcal{Y} \\ &= \operatorname{argmax}_{b \in \{0,1\}} \frac{\mathbb{P} \left(H_b^\ell \right) \cdot \mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell \right)}{\mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell \right)} \\ &= \operatorname{argmax}_{b \in \{0,1\}} \mathbb{P} \left(H_b^\ell \right) \cdot \mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell \right)\end{aligned}$$

- Where we take a Bayesian approach towards estimating the prior probabilities, $\mathbb{P} \left(H_b^\ell \right)$, and conditional probabilities, $\mathbb{P} \left(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell \right)$.

ML-KNN Algorithm: Overall Approach continued...

Definition: Let \vec{r}_t denote the real-valued vector with ℓ^{th} component:

$$\vec{r}_t(\ell) := \mathbb{P}\left(H_1^\ell\right).$$

ML-KNN Algorithm: Overall Approach continued...

Definition: Let \vec{r}_t denote the real-valued vector with ℓ^{th} component:

$$\vec{r}_t(\ell) := \mathbb{P} \left(H_1^\ell \right).$$

\Rightarrow Thus, given training data, \mathcal{X} , and a test instance, t , we wish to compute $[\vec{y}_t(\ell), \vec{r}_t(\ell)]$.

ML-KNN Algorithm: Computing the Prior Probabilities, $\widehat{\mathbb{P}}(H_b^\ell)$

We model $\mathbb{P}(H_1^\ell)$ with a $\text{Beta}(s, s)$ prior and $\text{Binomial}(m, \mathbb{P}(H_1^\ell))$ likelihood.
(When $s = 1$, $\text{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

ML-KNN Algorithm: Computing the Prior Probabilities,

$$\widehat{\mathbb{P}(\mathbf{H}_b^\ell)}$$

We model $\mathbb{P}(\mathbf{H}_1^\ell)$ with a $\text{Beta}(s, s)$ prior and $\text{Binomial}(m, \mathbb{P}(\mathbf{H}_1^\ell))$ likelihood. (When $s = 1$, $\text{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

\Rightarrow The posterior distribution for $\mathbb{P}(\mathbf{H}_1^\ell)$ is:

$$\text{Beta} \left(s + \sum_{i=1}^m \vec{y}_{x_i}(\ell), s + m - \sum_{i=1}^m \vec{y}_{x_i}(\ell) \right).$$

ML-KNN Algorithm: Computing the Prior Probabilities,

$$\widehat{\mathbb{P}(\mathbf{H}_b^\ell)}$$

We model $\mathbb{P}(\mathbf{H}_1^\ell)$ with a $\text{Beta}(s, s)$ prior and $\text{Binomial}(m, \mathbb{P}(\mathbf{H}_1^\ell))$ likelihood. (When $s = 1$, $\text{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

\Rightarrow The posterior distribution for $\mathbb{P}(\mathbf{H}_1^\ell)$ is:

$$\text{Beta} \left(s + \sum_{i=1}^m \vec{y}_{x_i}(\ell), s + m - \sum_{i=1}^m \vec{y}_{x_i}(\ell) \right).$$

\Rightarrow We will estimate $\mathbb{P}(\mathbf{H}_1^\ell)$ with the expectation of it's posterior Beta distribution:

$$\widehat{\mathbb{P}(\mathbf{H}_1^\ell)} := \frac{s + \sum_{i=1}^m \vec{y}_{x_i}(\ell)}{2s + m}$$

where m is the number of training instances.

ML-KNN Algorithm: Computing the Prior Probabilities, $\widehat{\mathbb{P}}(H_b^\ell)$

We model $\mathbb{P}(H_1^\ell)$ with a $\text{Beta}(s, s)$ prior and $\text{Binomial}(m, \mathbb{P}(H_1^\ell))$ likelihood. (When $s = 1$, $\text{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

\Rightarrow The posterior distribution for $\mathbb{P}(H_1^\ell)$ is:

$$\text{Beta} \left(s + \sum_{i=1}^m \vec{y}_{x_i}(\ell), s + m - \sum_{i=1}^m \vec{y}_{x_i}(\ell) \right).$$

\Rightarrow We will estimate $\mathbb{P}(H_1^\ell)$ with the expectation of it's posterior Beta distribution:

$$\widehat{\mathbb{P}}(H_1^\ell) := \frac{s + \sum_{i=1}^m \vec{y}_{x_i}(\ell)}{2s + m}$$

where m is the number of training instances.

\Rightarrow We estimate $\widehat{\mathbb{P}}(H_0^\ell) := 1 - \widehat{\mathbb{P}}(H_1^\ell)$.

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}}(E_j^\ell | H_b^\ell)$

Definition:

Model:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\mathbb{P}(\widehat{E_j^\ell} | \mathbf{H}_b^\ell)$

Definition:

- (i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.

Model:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\mathbb{P}(\widehat{E_j^\ell} | \mathbf{H}_b^\ell)$

Definition:

- (i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.
- (ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | H_b^\ell)}$

Definition:

- (i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.
- (ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

- * We let $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)} = (\mathbb{P}(E_1^\ell | H_1^\ell), \dots, \mathbb{P}(E_K^\ell | H_1^\ell))$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$)

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | H_b^\ell)}$

Definition:

- (i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.
- (ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

- * We let $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)} = (\mathbb{P}(E_1^\ell | H_1^\ell), \dots, \mathbb{P}(E_K^\ell | H_1^\ell))$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$)
- * We give $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ (and $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$) a $\text{Dirichlet}(K + 1, (s, \dots, s))$ prior distribution.

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | H_b^\ell)}$

Definition:

- (i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.
- (ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

- * We let $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)} = (\mathbb{P}(E_1^\ell | H_1^\ell), \dots, \mathbb{P}(E_K^\ell | H_1^\ell))$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$)
- * We give $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ (and $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$) a $\text{Dirichlet}(K + 1, (s, \dots, s))$ prior distribution.
- * We use a $\text{Multinomial}(K + 1, (\frac{c(0)}{m_1}, \dots, \frac{c(K)}{m_1}))$ likelihood, where $m_1 = \sum_{i=1}^m \vec{y}_{x_i}(\ell)$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$).

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | H_b^\ell)}$ continued...

\Rightarrow The posterior distribution for $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ is

$$\text{Dirichlet}\left(K + 1, \left(s + c(0), \dots, s + c(K)\right)\right)$$

(and analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$).

ML-KNN Algorithm: Computing the Conditional Probabilities, $\mathbb{P}(\widehat{E_j^\ell} | H_b^\ell)$ continued...

\Rightarrow The posterior distribution for $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ is

$$\text{Dirichlet}\left(K + 1, \left(s + c(0), \dots, s + c(K)\right)\right)$$

(and analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$).

\Rightarrow Given $j \in \{0, \dots, K\}$, we estimate $\mathbb{P}(E_j^\ell | H_1^\ell)$ and $\mathbb{P}(E_j^\ell | H_0^\ell)$ with the expectations of their posterior distributions:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\mathbb{P}(\widehat{E_j^\ell} | H_b^\ell)$ continued...

\Rightarrow The posterior distribution for $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ is

$$\text{Dirichlet}\left(K + 1, \left(s + c(0), \dots, s + c(K)\right)\right)$$

(and analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$).

\Rightarrow Given $j \in \{0, \dots, K\}$, we estimate $\mathbb{P}(E_j^\ell | H_1^\ell)$ and $\mathbb{P}(E_j^\ell | H_0^\ell)$ with the expectations of their posterior distributions:

$$\mathbb{P}(E_j^\ell | H_1^\ell) := \frac{(s + c(j))}{((K + 1)s + \sum_{n=0}^K c(n))}$$
$$\mathbb{P}(E_j^\ell | H_0^\ell) := \frac{(s + c'(j))}{((K + 1)s + \sum_{n=0}^K c'(n))}$$

Computing \vec{y}_t and \vec{r}_t

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \operatorname{argmax}_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \widehat{\mathbb{P}(E_{\vec{C}_T(\ell)}^\ell | H_b^\ell)} \right]$$

Computing \vec{y}_t and \vec{r}_t

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \operatorname{argmax}_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \widehat{\mathbb{P}(E_{\vec{C}_T(\ell)}^\ell | H_b^\ell)} \right]$$

AND

Computing \vec{y}_t and \vec{r}_t

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \operatorname{argmax}_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_b^\ell) \right]$$

AND

$$\begin{aligned} \widehat{\vec{r}_t(\ell)} &:= \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_1^\ell)}{\sum_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_b^\ell) \right]} \\ &= \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_1^\ell)}{\mathbb{P}(E_{\vec{C}_t(\ell)}^\ell)} \end{aligned}$$

Computing \vec{y}_t and \vec{r}_t

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \operatorname{argmax}_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_b^\ell) \right]$$

AND

$$\begin{aligned} \widehat{\vec{r}_t(\ell)} &:= \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_1^\ell)}{\sum_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_b^\ell) \right]} \\ &= \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell} | H_1^\ell)}{\mathbb{P}(E_{\vec{C}_t(\ell)}^\ell)} \end{aligned}$$

NOTE: The larger the value for s , the less importance assigned to the training data: As $s \rightarrow \infty$, $\vec{r}_t(\ell) \rightarrow \frac{1}{2}$.

Scikit-multilearn:

- The “**scikit-learn**” module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.

Scikit-multilearn:

- The “**scikit-learn**” module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.
- The “**scikit-multilearn**” module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

Scikit-multilearn:

- The “**scikit-learn**” module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.
- The “**scikit-multilearn**” module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

ML-KNN Classification in scikit-multilearn:

Scikit-multilearn Implementation

Scikit-multilearn:

- The “**scikit-learn**” module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.
- The “**scikit-multilearn**” module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

ML-KNN Classification in scikit-multilearn:

→ MLkNN() from the scikit-multilearn module can be used to instantiate a ML-KNN object.

```
from skmultilearn.adapt import MLkNN

classifier = MLkNN(k=3)

# train
classifier.fit(X_train, y_train)

# predict
predictions = classifier.predict(X_test)
```

Scikit-multilearn Implementation

Scikit-multilearn:

- The “**scikit-learn**” module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.
- The “**scikit-multilearn**” module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

ML-KNN Classification in scikit-multilearn:

- MLkNN() from the scikit-multilearn module can be used to instantiate a ML-KNN object.
- “MLkNN” class methods like “fit()” and “predict()” mirror those for standard scikit-learn objects.

```
from skmultilearn.adapt import MLkNN

classifier = MLkNN(k=3)

# train
classifier.fit(X_train, y_train)

# predict
predictions = classifier.predict(X_test)
```


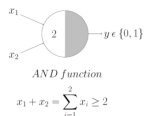
BP-MLL Approach

A Brief History

A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.



AND Function



A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.
- Since most subsequent work in the following two decades centered around single layer networks, the power of neural networks was restricted to linearly separable problems.

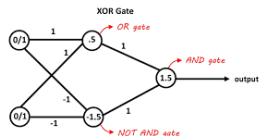

AND Function



A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.
- Since most subsequent work in the following two decades centered around single layer networks, the power of neural networks was restricted to linearly separable problems. This excluded the possibility of learning even simple functions like XOR, which required a second layer.




AND Function



A Brief History


- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.
- Since most subsequent work in the following two decades centered around single layer networks, the power of neural networks was restricted to linearly separable problems. This excluded the possibility of learning even simple functions like XOR, which required a second layer.
- In the early 1980s, research on neural networks resurged largely due to successful learning algorithms for multi-layer neural networks and are used today for various tasks such as computer vision, associative memory, representation learning, NLP, etc..

AND Function



Introduction to Feed-forward Networks: Perceptron Model

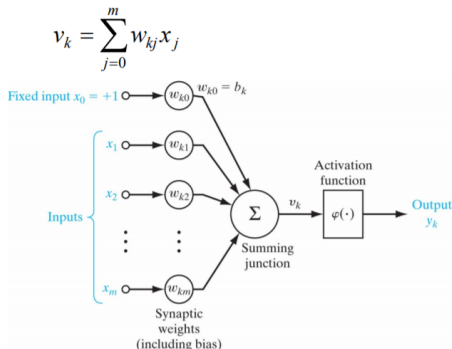
The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.



Introduction to Feed-forward Networks: Perceptron Model

The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.

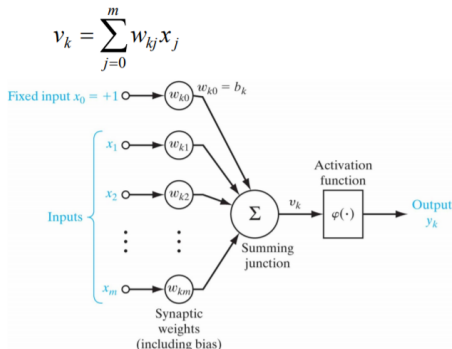
- Neurons in the first layer represent components of the input vectors.



Introduction to Feed-forward Networks: Perceptron Model

The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.

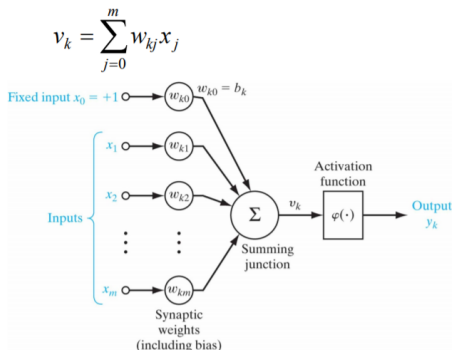
- Neurons in the first layer represent components of the input vectors.
- The output of the neuron in the next layer is determined by applying a non-linear “activation function” to a linear combination of the input components, plus a bias.



Introduction to Feed-forward Networks: Perceptron Model

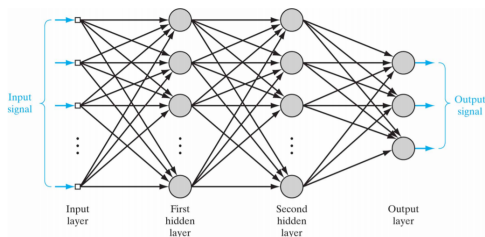
The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.

- Neurons in the first layer represent components of the input vectors.
- The output of the neuron in the next layer is determined by applying a non-linear “activation function” to a linear combination of the input components, plus a bias.
- Rosenblatt’s Perceptron model uses a step function non-linearity, but other common activation functions include the sigmoid function ($\sigma()$), $\tanh()$, ReLU, Leaky ReLU, softmax, etc..



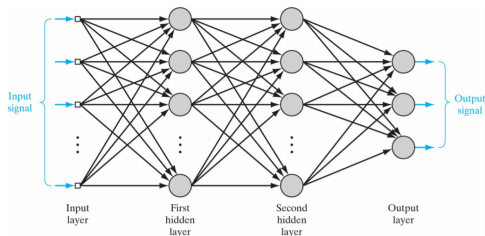
Multilayer Networks & Training

- Adding additional layers and units (like in the network pictured below) significantly expands the class of discrimination problems a network can learn.



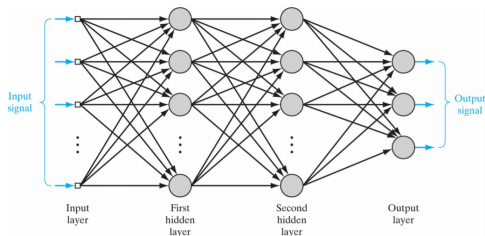
Multilayer Networks & Training

- Adding additional layers and units (like in the network pictured below) significantly expands the class of discrimination problems a network can learn.
- “Online” training involves evaluating an instance, and updating weights using gradient descent.



Multilayer Networks & Training

- Adding additional layers and units (like in the network pictured below) significantly expands the class of discrimination problems a network can learn.
- “Online” training involves evaluating an instance, and updating weights using gradient descent.
- For multi-label learning with Q instances, networks will have Q output layers, each with a $\tanh()$ activation.



Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Naive Approach (“BasicBP”):

- Standard MSE:

$$E = \sum_{i=1}^m E_i = \sum_{i=1}^m \sum_{j=1}^q (c_j^i - d_j^i)^2$$

where $c_j^i = c_j(x_i)$ is the output of the network on x_i on the j^{th} class.

Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Naive Approach (“BasicBP”):

- Standard MSE:

$$E = \sum_{i=1}^m E_i = \sum_{i=1}^m \sum_{j=1}^q (c_j^i - d_j^i)^2$$

where $c_j^i = c_j(x_i)$ is the output of the network on x_i on the j^{th} class.

Novel Approach (“BP-MLL”):

- BP-MLL Cost function:

$$E = \sum_{i=1}^m E_i = \sum_{i=1}^m \frac{1}{|Y_i| |\bar{Y}_i|} \sum_{(k,l) \in Y_i \times \bar{Y}_i} \exp(-(c_k^i - c_l^i))$$

so that the i^{th} error term is severely penalized if c_k^i is much smaller than c_l^i .

Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Naive Approach (“BasicBP”):

- Standard MSE:

$$E = \sum_{i=1}^m E_i = \sum_{i=1}^m \sum_{j=1}^q (c_j^i - d_j^i)^2$$

where $c_j^i = c_j(x_i)$ is the output of the network on x_i on the j^{th} class.

Novel Approach (“BP-MLL”):

- BP-MLL Cost function:

$$E = \sum_{i=1}^m E_i = \sum_{i=1}^m \frac{1}{|Y_i| |\bar{Y}_i|} \sum_{(k,l) \in Y_i \times \bar{Y}_i} \exp(-(c_k^i - c_l^i))$$

so that the i^{th} error term is severely penalized if c_k^i is much smaller than c_l^i .

- Back-propagation for training is derived just as in the standard (MSE) case (details omitted here, but can be found in Zhang and Zhou's paper).

Deep Learning APIs in Python: TensorFlow/Keras

TensorFlow: an open source python library for numerical computation and large-scale machine learning, created by the Google Brain team.

- One of the most widely used APIs for deep learning, along with PyTorch and Keras.
- Later versions of TensorFlow began incorporating the Keras API, since users found its high-level design to be simpler.

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation=tf.nn.relu),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TensorFlow Implementation of BP-MLL

BP-MLL in TensorFlow:

- Data Scientist, Lukas Huwald, published an implementation of the bp-mlm cost function for the TensorFlow API as part of the module “bpmll”.
- After installation, the bp-mlm loss function can be utilized just as any other TensorFlow loss function.

```
# create simple mlp
model = Sequential()
model.add(Dense(128, input_dim=dim_no, activation='relu', kernel_initializer='glorot_uniform'))
model.add(Dense(64, activation='relu', kernel_initializer='glorot_uniform'))
model.add(Dense(class_no, activation='sigmoid', kernel_initializer='glorot_uniform'))
model.compile(loss=bp_mll_loss, optimizer='adagrad', metrics=[])

# train a few epochs
model.fit(X_train, Y_train, epochs=100)
```


Potential Cause for Caution

Both Zhang and Zhou's original paper as well as Nam et al. [2014] provide reasons for potential caution when considering application of the BP-MLL loss function:

Potential Cause for Caution

Both Zhang and Zhou's original paper as well as Nam et al. [2014] provide reasons for potential caution when considering application of the BP-MLL loss function:

- **Computational Inefficiency:** Because the BP-MLL loss function involves pairwise comparisons, obtaining error terms is more expensive than utilizing cross-entropy or MSE loss.

Potential Cause for Caution

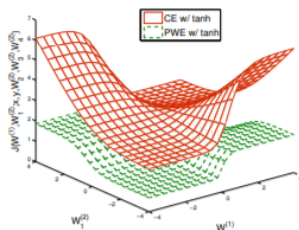
Both Zhang and Zhou's original paper as well as Nam et al. [2014] provide reasons for potential caution when considering application of the BP-MLL loss function:

- **Computational Inefficiency:** Because the BP-MLL loss function involves pairwise comparisons, obtaining error terms is more expensive than utilizing cross-entropy or MSE loss. This scales poorly with the number of labels, and can lead to significantly larger training times.

Potential Cause for Caution

Both Zhang and Zhou's original paper as well as Nam et al. [2014] provide reasons for potential caution when considering application of the BP-MLL loss function:

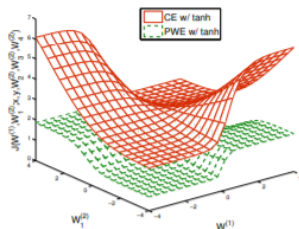
- **Computational Inefficiency:** Because the BP-MLL loss function involves pairwise comparisons, obtaining error terms is more expensive than utilizing cross-entropy or MSE loss. This scales poorly with the number of labels, and can lead to significantly larger training times.
- **Objective Function Surface:** The surface for the BP-MLL loss has plateaus in which gradient descent can be very slow in comparison with the cross-entropy, for example (Nam et al. [2014]).



Potential Cause for Caution

Both Zhang and Zhou's original paper as well as Nam et al. [2014] provide reasons for potential caution when considering application of the BP-MLL loss function:

- **Computational Inefficiency:** Because the BP-MLL loss function involves pairwise comparisons, obtaining error terms is more expensive than utilizing cross-entropy or MSE loss. This scales poorly with the number of labels, and can lead to significantly larger training times.
- **Objective Function Surface:** The surface for the BP-MLL loss has plateaus in which gradient descent can be very slow in comparison with the cross-entropy, for example (Nam et al. [2014]).



- **Better Generalization:** When compared against a “standard” feed forward network with dropout regularization, adaptive learning rates and ReLU hidden layer activations, the test-set performance of BP-MLL is inferior on benchmark datasets for large scale text categorization (Nam et al. [2014]).

References (Original Method Papers)

- Jinseok Nam, Jungi Kim, Eneldo Loza Mencía, Iryna Gurevych, and Johannes Fürnkranz. Large-scale multi-label text classification — revisiting neural networks. In Toon Calders, Floriana Esposito, Eyke Hüllermeier, and Rosa Meo, editors, *Machine Learning and Knowledge Discovery in Databases*, pages 437–452, Berlin, Heidelberg, 2014. Springer Berlin Heidelberg. ISBN 978-3-662-44851-9.
- Min-Ling Zhang and Zhi-Hua Zhou. Ml-knn: A lazy learning approach to multi-label learning. *Pattern Recognition*, 40(7):2038–2048, 2007. doi: 10.1016/j.patcog.2006.12.019.
- Min-Ling Zhang and Zhi-Hua Zhou. Multilabel neural networks with applications to functional genomics and text categorization. *IEEE Transactions on Knowledge and Data Engineering*, 18(10):1338–1351, 2006. doi: doi:10.1109/TKDE.2006.162.