An Introduction to Multi-label Learning
(ML-KNN & BP-MLL)

Bobby Lumpkin

Overview

Introduction to Multi-label Learning

Multi-label Learning

What is it?

Multi-label Learning

What is it?

$\rightarrow$ Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

Multi-label Learning

What is it?

$\rightarrow$ Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

$\rightarrow$ Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.

Multi-label Learning

What is it?

$\rightarrow$ Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

$\rightarrow$ Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.

* **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: "News", "Health" and "Government".

Multi-label Learning

What is it?

$\rightarrow$ Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

$\rightarrow$ Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.

* **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: "News", "Health" and "Government".

Why use novel approaches?

Multi-label Learning

What is it?

$\rightarrow$ Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

$\rightarrow$ Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.

* **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: "News", "Health" and "Government".

Why use novel approaches?

$\rightarrow$ Naive Approach: Train a sequence of independent binary classifiers (one per category)

Multi-label Learning

What is it?

$\rightarrow$ Multi-label learning (mll) is a form of classification where each instance may be associated with more than one label.

$\rightarrow$ Plenty of tasks, such as text categorization, functional genomics, and supervised product recommendation fit naturally into the mll paradigm.

* **EX:** A news article discussing White House Covid press briefings might belong to each of the categories: "News", "Health" and "Government".

Why use novel approaches?

$\rightarrow$ Naive Approach: Train a sequence of independent binary classifiers (one per category)

$\rightarrow$ Doesn't capitalize on the information in the correlations between the different labels of each instance.

Multi-label Paradigm: Definitions & Notation

- Let χ denote the domain of instances and $\mathcal{Y} = \{1, ..., Q\}$ be the finite set of labels.
- Given $x \in \chi$ and its associated $Y \subseteq \mathcal{Y}$, let $\vec{y}_x$ be the category vector for x such that (for all $\ell \in \mathcal{Y}$) $\vec{y}_x(\ell) = 1$ if $\ell \in Y$. Otherwise, $\vec{y}_x(\ell) = 0$.

ML-KNN Approach

ML-KNN Algorithm: More Notation

Notation:

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x, identified in the training set.

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x, identified in the training set.
- Let $\vec{C}_x(\ell) = \sum_{a \in N(x)} \vec{y}_a(\ell)$ ($\ell \in \mathcal{Y}$) define a membership counting vector.

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x, identified in the training set.
- Let $\vec{C}_x(\ell) = \sum_{a \in N(x)} \vec{y}_a(\ell)$ $(\ell \in \mathcal{Y})$ define a membership counting vector.
- Let H_0^ℓ denote the event that test instance t does not have a label ℓ and let H_1^ℓ denote the event that it does have label ℓ.

ML-KNN Algorithm: More Notation

Notation:

- Let $N(x)$ denote the set of K nearest neighbors of x, identified in the training set.
- Let $\vec{C}_x(\ell) = \sum_{a \in N(x)} \vec{y}_a(\ell)$ ($\ell \in \mathcal{Y}$) define a membership counting vector.
- Let H_0^ℓ denote the event that test instance t does not have a label ℓ and let H_1^ℓ denote the event that it does have label ℓ.
- Let E_j^ℓ ($j \in \{1, ..., K\}$) denote the event that, among the K nearest neighbors of t, there are exactly j instances which have label ℓ.

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

- Given a test instance, t, $\vec{Y}_t$ is determined using the MAP estimate:

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

- Given a test instance, t, $\vec{Y}_t$ is determined using the MAP estimate:

$$
\vec{y}_t(\ell) = \operatorname*{argmax}_{b \in \{0,1\}} \mathbb{P}\left(\mathrm{H}_b^\ell | E_{\vec{C}_t(\ell)}^\ell\right), \quad \ell \in \mathcal{Y}
$$

$$
= \operatorname*{argmax}_{b \in \{0,1\}} \frac{\mathbb{P}\left(\mathrm{H}_b^\ell\right) \cdot \mathbb{P}\left(E_{\vec{C}_{t(\ell)}}^\ell | \mathrm{H}_b^\ell\right)}{\mathbb{P}\left(E_{\vec{C}_t(\ell)}^\ell\right)}
$$

$$
= \operatorname*{argmax}_{b \in \{0,1\}} \mathbb{P}\left(\mathrm{H}_b^\ell\right) \cdot \mathbb{P}\left(E_{\vec{C}_{t(\ell)}}^\ell | \mathrm{H}_b^\ell\right)
$$

ML-KNN Algorithm: Overall Approach

Overall Approach: This ML-KNN algorithm takes a parametric, Bayesian approach towards estimating the Bayes Optimal Classifier. As with the single-label algorithm, it does this using the K nearest neighbors of an instance. Namely...

- Given a test instance, t, $\vec{Y}_t$ is determined using the MAP estimate:

$$
\begin{aligned}
\vec{y}_t(\ell) &= \operatorname*{argmax}_{b \in \{0,1\}} \mathbb{P}\left(\mathrm{H}_b^\ell | E_{\vec{C}_t(\ell)}^\ell\right), \quad \ell \in \mathcal{Y} \\
&= \operatorname*{argmax}_{b \in \{0,1\}} \frac{\mathbb{P}\left(\mathrm{H}_b^\ell\right) \cdot \mathbb{P}\left(E_{\vec{C}_{t(\ell)}}^\ell | \mathrm{H}_b^\ell\right)}{\mathbb{P}\left(E_{\vec{C}_t(\ell)}^\ell\right)} \\
&= \operatorname*{argmax}_{b \in \{0,1\}} \mathbb{P}\left(\mathrm{H}_b^\ell\right) \cdot \mathbb{P}\left(E_{\vec{C}_{t(\ell)}}^\ell | \mathrm{H}_b^\ell\right)
\end{aligned}
$$

- Where we take a Bayesian approach towards estimating the prior probabilities, $\mathbb{P}\left(\mathrm{H}_b^\ell\right)$, and conditional probabilities, $\mathbb{P}\left(E_{\vec{C}_{t(\ell)}}^\ell | \mathrm{H}_b^\ell\right)$.

Definition: Let $\vec{r}_t$ denote the real-valued vector with ℓ^{th} component:

$$\vec{r}_t(\ell) := \mathbb{P}\left(\mathrm{H}_1^\ell\right).$$

ML-KNN Algorithm: Overall Approach continued...

Definition: Let $\vec{r}_t$ denote the real-valued vector with ℓ^{th} component:

$$\vec{r}_t(\ell) \coloneqq \mathbb{P}\left(H_1^\ell\right).$$

$\Rightarrow$ Thus, given training data, $\mathcal{X}$, and a test instance, t, we wish to compute $[\vec{y}_t(\ell), \vec{r}_t(\ell)]$.

ML-KNN Algorithm: Computing the Prior Probabilities, $\widehat{\mathbb{P}(H_b^\ell)}$

We model $\mathbb{P}(H_1^\ell)$ with a $\mathrm{Beta}(s, s)$ prior and $\mathrm{Binomial}(m, \ \mathbb{P}(H_1^\ell))$ likelihood.
(When $s = 1$, $\mathrm{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

ML-KNN Algorithm: Computing the Prior Probabilities, $\widehat{\mathbb{P}(H_b^\ell)}$

We model $\mathbb{P}(H_1^\ell)$ with a $\mathrm{Beta}(s, s)$ prior and $\mathrm{Binomial}(m, \ \mathbb{P}(H_1^\ell))$ likelihood. (When $s = 1$, $\mathrm{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

$\Rightarrow$ The posterior distribution for $\mathbb{P}(H_1^\ell)$ is:

$$\mathrm{Beta}\left(s + \sum_{i=1}^m \vec{y}_{x_i}(\ell), \ s + m - \sum_{i=1}^m \vec{y}_{x_i}(\ell)\right).$$

ML-KNN Algorithm: Computing the Prior Probabilities, $\widehat{\mathbb{P}(H_b^\ell)}$

We model $\mathbb{P}(H_1^\ell)$ with a $\mathrm{Beta}(s, s)$ prior and $\mathrm{Binomial}(m, \ \mathbb{P}(H_1^\ell))$ likelihood. (When $s = 1$, $\mathrm{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

$\Rightarrow$ The posterior distribution for $\mathbb{P}(H_1^\ell)$ is:

$$\mathrm{Beta}\left(s + \sum_{i=1}^{m} \vec{y}_{x_i}(\ell), \ s + m - \sum_{i=1}^{m} \vec{y}_{x_i}(\ell) \right).$$

$\Rightarrow$ We will estimate $\mathbb{P}(H_1^\ell)$ with the expectation of it's posterior Beta distribution:

$$\widehat{\mathbb{P}(H_1^\ell)} := \frac{s + \sum_{i=1}^{m} \vec{y}_{x_i}(\ell)}{2s + m}$$

where m is the number of training instances.

ML-KNN Algorithm: Computing the Prior Probabilities, $\widehat{\mathbb{P}(H_b^\ell)}$

We model $\mathbb{P}(H_1^\ell)$ with a $\mathrm{Beta}(s, s)$ prior and $\mathrm{Binomial}(m, \; \mathbb{P}(H_1^\ell))$ likelihood.
(When $s = 1$, $\mathrm{Beta}(s, s)$ reduces to the uniform distribution on $[0, 1]$.)

$\Rightarrow$ The posterior distribution for $\mathbb{P}(H_1^\ell)$ is:

$$\mathrm{Beta}\left(s + \sum_{i=1}^m \vec{y}_{x_i}(\ell), \; s + m - \sum_{i=1}^m \vec{y}_{x_i}(\ell)\right).$$

$\Rightarrow$ We will estimate $\mathbb{P}(H_1^\ell)$ with the expectation of it's posterior Beta distribution:

$$\widehat{\mathbb{P}(H_1^\ell)} := \frac{s + \sum_{i=1}^m \vec{y}_{x_i}(\ell)}{2s + m}$$

where m is the number of training instances.

$\Rightarrow$ We estimate $\widehat{\mathbb{P}(H_0^\ell)} := 1 - \widehat{\mathbb{P}(H_1^\ell)}$.

Definition:

Model:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | H_b^\ell)}$

Definition:

(i) Let c be a vector of length $K+1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.

Model:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | \mathrm{H}_b^\ell)}$

Definition:

(i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.

(ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | \mathrm{H}_b^\ell)}$

Definition:

(i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.

(ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

* We let $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_1^\ell)} = (\mathbb{P}(E_1^\ell | \mathrm{H}_1^\ell), ..., \mathbb{P}(E_K^\ell | \mathrm{H}_1^\ell))$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_0^\ell)}$)

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | \mathrm{H}_b^\ell)}$

Definition:

(i) Let c be a vector of length $K+1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.

(ii) Similarly, let c' be a vector of length $K+1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

* We let $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_1^\ell)} = (\mathbb{P}(E_1^\ell | \mathrm{H}_1^\ell), ..., \mathbb{P}(E_K^\ell | \mathrm{H}_1^\ell))$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_0^\ell)}$)

* We give $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_1^\ell)}$ (and $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_0^\ell)}$) a $\mathrm{Dirichlet}(K+1, (s, ..., s))$ prior distribution.

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | \mathrm{H}_b^\ell)}$

Definition:

(i) Let c be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 1$.

(ii) Similarly, let c' be a vector of length $K + 1$, where $c(j) =$ the number of training instances where $\vec{C}_{x_i}(\ell) = j$ when $\vec{y}_{x_i}(\ell) = 0$.

Model:

* We let $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_1^\ell)} = (\mathbb{P}(E_1^\ell | \mathrm{H}_1^\ell), ..., \mathbb{P}(E_K^\ell | \mathrm{H}_1^\ell))$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_0^\ell)}$)

* We give $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_1^\ell)}$ (and $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_0^\ell)}$) a $\mathrm{Dirichlet}(K + 1, (s, ..., s))$ prior distribution.

* We use a $\mathrm{Multinomial}(K + 1, (\frac{c(0)}{m_1}, ..., \frac{c(K)}{m_1}))$ likelihood, where $m_1 = \sum_{i=1}^m \vec{y}_{x_i}(\ell)$ (analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | \mathrm{H}_0^\ell)}$).

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | \mathrm{H}_b^\ell)}$ continued...

$\Rightarrow$ The posterior distribution for $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ is

$$\mathrm{Dirichlet}\left(K + 1, \left(s + c(0), ..., s + c(K) \right) \right)$$

(and analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$.

ML-KNN Algorithm: Computing the Conditional Probabilities, $\widehat{\mathbb{P}(E_j^\ell | \mathrm{H}_b^\ell)}$ continued...

$\Rightarrow$ The posterior distribution for $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ is

$$\mathrm{Dirichlet}\left(K + 1, \left(s + c(0), ..., s + c(K) \right) \right)$$

(and analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$).

$\Rightarrow$ Given $j \in \{0, ..., K\}$, we estimate $\mathbb{P}(E_j^\ell | H_1^\ell)$ and $\mathbb{P}(E_j^\ell | H_0^\ell)$ with the expectations of their posterior distributions:

$\Rightarrow$ The posterior distribution for $\overrightarrow{\mathbb{P}(E_j^\ell | H_1^\ell)}$ is

$$\mathrm{Dirichlet}\left(K + 1, \left(s + c(0), ..., s + c(K) \right) \right)$$

(and analogously for $\overrightarrow{\mathbb{P}(E_j^\ell | H_0^\ell)}$).

$\Rightarrow$ Given $j \in \{0, ..., K\}$, we estimate $\mathbb{P}(E_j^\ell | H_1^\ell)$ and $\mathbb{P}(E_j^\ell | H_0^\ell)$ with the expectations of their posterior distributions:

$$\mathbb{P}(E_j^\ell | H_1^\ell) := \frac{(s + c(j))}{((K + 1)s + \sum_{n=0}^{K} c(n))}$$

$$\mathbb{P}(E_j^\ell | H_0^\ell) := \frac{(s + c'(j))}{((K + 1)s + \sum_{n=0}^{K} c'(n))}$$

Computing $\vec{y}_t$ and $\vec{r}_t$

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \operatorname*{argmax}_{b \in \{0,1\}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \widehat{\mathbb{P}(E_{\bar{C}_T(\ell)}^\ell | H_b^\ell)} \right]$$

Computing $\vec{y}_t$ and $\vec{r}_t$

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \underset{b \in \{0,1\}}{\operatorname{argmax}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \widehat{\mathbb{P}(E_{\hat{C}_T(\ell)}^\ell | H_b^\ell)} \right]$$

AND

Computing $\vec{y}_t$ and $\vec{r}_t$

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \underset{b \in \{0,1\}}{\operatorname{argmax}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell | H_b^\ell}) \right]$$

AND

$$\widehat{\vec{r}_t(\ell)} := \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell | H_1^\ell})}{\sum_{b \in \{0,1\}} \left[\mathbb{P}(H_b^\ell) \cdot \mathbb{P}(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell) \right]}$$

$$= \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell | H_1^\ell})}{\mathbb{P}(E_{\vec{C}_t(\ell)}^\ell)}$$

Computing $\vec{y}_t$ and $\vec{r}_t$

Using our previous derivation:

$$\widehat{\vec{y}_t(\ell)} := \underset{b \in \{0,1\}}{\operatorname{argmax}} \left[\widehat{\mathbb{P}(H_b^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell | H_b^\ell}) \right]$$

AND

$$\widehat{\vec{r}_t(\ell)} := \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell | H_1^\ell})}{\sum_{b \in \{0,1\}} \left[\mathbb{P}(H_b^\ell) \cdot \mathbb{P}(E_{\vec{C}_t(\ell)}^\ell | H_b^\ell) \right]}$$

$$= \frac{\widehat{\mathbb{P}(H_1^\ell)} \cdot \mathbb{P}(\widehat{E_{\vec{C}_T(\ell)}^\ell | H_1^\ell})}{\mathbb{P}(E_{\vec{C}_t(\ell)}^\ell)}$$

NOTE: The larger the value for s, the less importance assigned to the training data: As $s \to \infty$, $\vec{r}_t(\ell) \to \frac{1}{2}$.

Scikit-multilearn Implementation

Scikit-multilearn:

- The **"scikit-learn"** module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.

Scikit-multilearn Implementation

Scikit-multilearn:

- The **"scikit-learn"** module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.

- The **"scikit-multilearn"** module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

Scikit-multilearn Implementation

Scikit-multilearn:

- The **"scikit-learn"** module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.

- The **"scikit-multilearn"** module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

ML-KNN Classification in scikit-multilearn:

Scikit-multilearn Implementation

Scikit-multilearn:

- The **"scikit-learn"** module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.

- The **"scikit-multilearn"** module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

ML-KNN Classification in scikit-multilearn:

$\rightarrow$ MLkNN() from the scikit-multilearn module can be used to instantiate a ML-KNN object.

```python
from skmultilearn.adapt import MLkNN

classifier = MLkNN(k=3)

# train
classifier.fit(X_train, y_train)

# predict
predictions = classifier.predict(X_test)
```

Scikit-multilearn Implementation

Scikit-multilearn:

- The **"scikit-learn"** module is a free and widely used software machine learning library for Python, including many popular regression, classification and unsupervised learning algorithms.

- The **"scikit-multilearn"** module is a library for multi-label classification that is built on top of the scikit-learn ecosystem.

ML-KNN Classification in scikit-multilearn:

$\rightarrow$ MLkNN() from the scikit-multilearn module can be used to instantiate a ML-KNN object.

$\rightarrow$ "MLkNN" class methods like "fit()" and "predict()" mirror those for standard scikit-learn objects.

```python
from skmultilearn.adapt import MLkNN

classifier = MLkNN(k=3)

# train
classifier.fit(X_train, y_train)

# predict
predictions = classifier.predict(X_test)
```

BP-MLL Approach

A Brief History

A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.
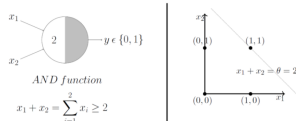
AND Function

$$AND\ function$$

$$x_1 + x_2 = \sum_{i=1}^{2} x_i \geq 2$$

A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.

- Since most subsequent work in the following two decades centered around single layer networks, the power of neural networks was restricted to linearly separable problems.

AND Function

$AND\ function$

$$x_1 + x_2 = \sum_{i=1}^{2} x_i \geq 2$$

$$x_1 + x_2 = \theta = 2$$

A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.

- Since most subsequent work in the following two decades centered around single layer networks, the power of neural networks was restricted to linearly separable problems. This excluded the possibility of learning even simple functions like XOR, which required a second layer.

AND Function

$$AND\ function$$

$$x_1 + x_2 = \sum_{i=1}^{2} x_i \geq 2$$

XOR Gate

A Brief History

- Inspired by biological nervous systems, neural networks date back to the first half of the 20th century with works such as those by McCulloch and Pitts, which could model simple logical operations.

- Since most subsequent work in the following two decades centered around single layer networks, the power of neural networks was restricted to linearly separable problems. This excluded the possibility of learning even simple functions like XOR, which required a second layer.

- In the early 1980s, research on neural networks resurged largely due to successful learning algorithms for multi-layer neural networks and are used today for various tasks such as computer vision, associative memory, representation learning, NLP, etc..

AND Function

$$x_1 + x_2 = \sum_{i=1}^{2} x_i \geq 2$$

XOR Gate

Hidden layer

Visible layer
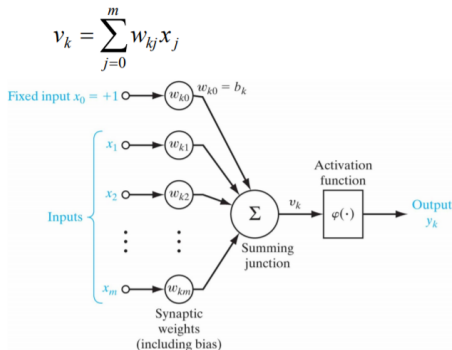
Introduction to Feed-forward Networks: Perceptron Model

The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.

$$v_k = \sum_{j=0}^{m} w_{kj} x_j$$

Introduction to Feed-forward Networks: Perceptron Model

The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.
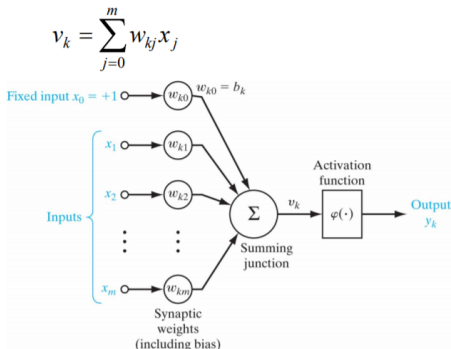
- Neurons in the first layer represent components of the input vectors.

$$v_k = \sum_{j=0}^{m} w_{kj} x_j$$

Introduction to Feed-forward Networks: Perceptron Model

The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.
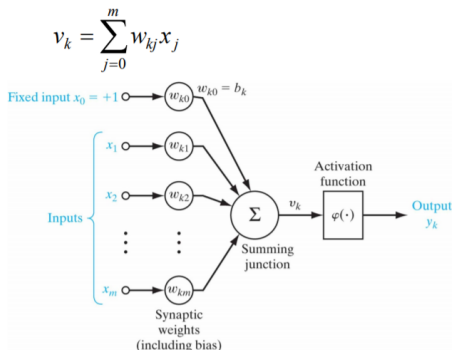
- Neurons in the first layer represent components of the input vectors.

- The output of the neuron in the next layer is determined by applying a non-linear "activation function" to a linear combination of the input components, plus a bias.

$$v_k = \sum_{j=0}^{m} w_{kj} x_j$$

## Introduction to Feed-forward Networks: Perceptron Model

The perceptron model (the building block for feed-forward networks) can be viewed as a connected, directed, loop-free graph like the one below.

- Neurons in the first layer represent components of the input vectors.

- The output of the neuron in the next layer is determined by applying a non-linear "activation function" to a linear combination of the input components, plus a bias.

- Rosenblatt's Perceptron model uses a step function non-linearity, but other common activation functions include the sigmoid function ($\sigma()$), tanh(), ReLU, Leaky ReLU, softmax, etc..

$$v_k = \sum_{j=0}^{m} w_{kj} x_j$$

Fixed input $x_0 = +1$ ○——○ w_{k0} $w_{k0} = b_k$

Inputs
- x_1 ○——○ w_{k1}
- x_2 ○——○ w_{k2}
- ⋮
- x_m ○——○ w_{km}

Σ v_k $\varphi(\cdot)$ → Output y_k

Summing junction

Activation function
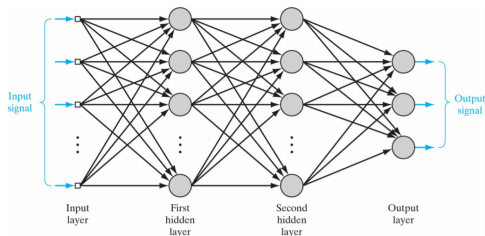
Synaptic weights (including bias)

Multilayer Networks & Training

- Adding additional layers and units (like in the network pictured below) significantly expands the class of discrimination problems a network can learn.
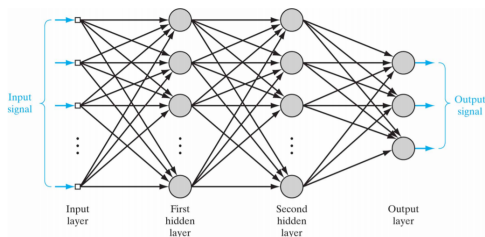
Multilayer Networks & Training

- Adding additional layers and units (like in the network pictured below) significantly expands the class of discrimination problems a network can learn.
- "Online" training involves evaluating an instance, and updating weights using gradient descent.

Multilayer Networks & Training

- Adding additional layers and units (like in the network pictured below) significantly expands the class of discrimination problems a network can learn.
- "Online" training involves evaluating an instance, and updating weights using gradient descent.
- For multi-label learning with Q instances, networks will have Q output layers, each with a tanh() activation.
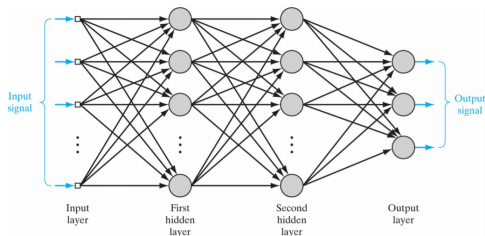
Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.
Naive Approach ("BasicBP"):

- Standard MSE:

$$E = \sum_{i=1}^{m} E_i = \sum_{i=1}^{m} \sum_{j=1}^{q} (c_j^i - d_j^i)^2$$

where $c_j^i = c_j(x_i)$ is the output of the network on x_i on the j^{th} class.

Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Naive Approach ("BasicBP"):

- Standard MSE:

$$E = \sum_{i=1}^{m} E_i = \sum_{i=1}^{m} \sum_{j=1}^{q} (c_j^i - d_j^i)^2$$

where $c_j^i = c_j(x_i)$ is the output of the network on x_i on the j^{th} class.

Novel Approach ("BP-MLL"):

- BP-MLL Cost function:

$$E = \sum_{i=1}^{m} E_i = \sum_{i=1}^{m} \frac{1}{|Y_i||\overline{Y}_i|} \sum_{(k,l) \in Y_i \times \overline{Y}_i} \exp(-(c_k^i - c_l^i))$$

so that the i^{th} error term is severely penalized if c_k^i is much smaller than c_l^i.

Designing a Cost Function

As for standard networks, BP-MLL uses gradient descent & back propagation for learning. The novelty of the approach is in the design of the cost function.

Naive Approach ("BasicBP"):

- Standard MSE:

$$E = \sum_{i=1}^{m} E_i = \sum_{i=1}^{m} \sum_{j=1}^{q} (c_j^i - d_j^i)^2$$

where $c_j^i = c_j(x_i)$ is the output of the network on x_i on the j^{th} class.

Novel Approach ("BP-MLL"):

- BP-MLL Cost function:

$$E = \sum_{i=1}^{m} E_i = \sum_{i=1}^{m} \frac{1}{|Y_i||\overline{Y}_i|} \sum_{(k,l) \in Y_i \times \overline{Y}_i} \exp(-(c_k^i - c_l^i))$$

so that the i^{th} error term is severely penalized if c_k^i is much smaller than c_l^i.

- Back-propagation for training is derived just as in the standard (MSE) case (details omitted here, but can be found in Zhang and Zhou's paper).

Deep Learning APIs in Python: TensorFlow/Keras

TensorFlow: an open source python library for numerical computation and large-scale machine learning, created by the Google Brain team.

$\rightarrow$ One of the most widely used APIs for deep learning, along with PyTorch and Keras.

$\rightarrow$ Later versions of TensorFlow began incorporating the Keras API, since users found its high-level design to be simpler.

```python
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train),(x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
  tf.keras.layers.Flatten(),
  tf.keras.layers.Dense(512, activation=tf.nn.relu),
  tf.keras.layers.Dropout(0.2),
  tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TensorFlow Implementation of BP-MLL

BP-MLL in TensorFlow:

$\rightarrow$ Data Scientist, Lukas Huwald, published an implementation of the bp-mll cost function for the TensorFlow API as part of the module "bpmll".

$\rightarrow$ After installation, the bp-mll loss function can be utilized just as any other TensorFlow loss function.

```
# create simple mlp
model = Sequential()
model.add(Dense(128, input_dim=dim_no, activation='relu', kernel_initializer='glorot_uniform'))
model.add(Dense(64, activation='relu', kernel_initializer='glorot_uniform'))
model.add(Dense(class_no, activation='sigmoid', kernel_initializer='glorot_uniform'))
model.compile(loss=bp_mll_loss, optimizer='adagrad', metrics=[])

# train a few epochs
model.fit(X_train, Y_train, epochs=100)
```

References (Original Method Papers)

Min-Ling Zhang and Zhi-Hua Zhou. Ml-knn: A lazy learning approach to multi-label learning. *Pattern Recognition*, 40(7):2038–2048, 2007. doi: 10.1016/j.patcog.2006.12.019.

Min-Ling Zhang and Zhi-Hua Zhou. Multilabel neural networks with applications to functional genomics and text categorization. *IEEE Transactions on Knowledge and Data Engineering*, 18(10):1338–1351, 2006. doi: doi:10.1109/TKDE.2006.162.